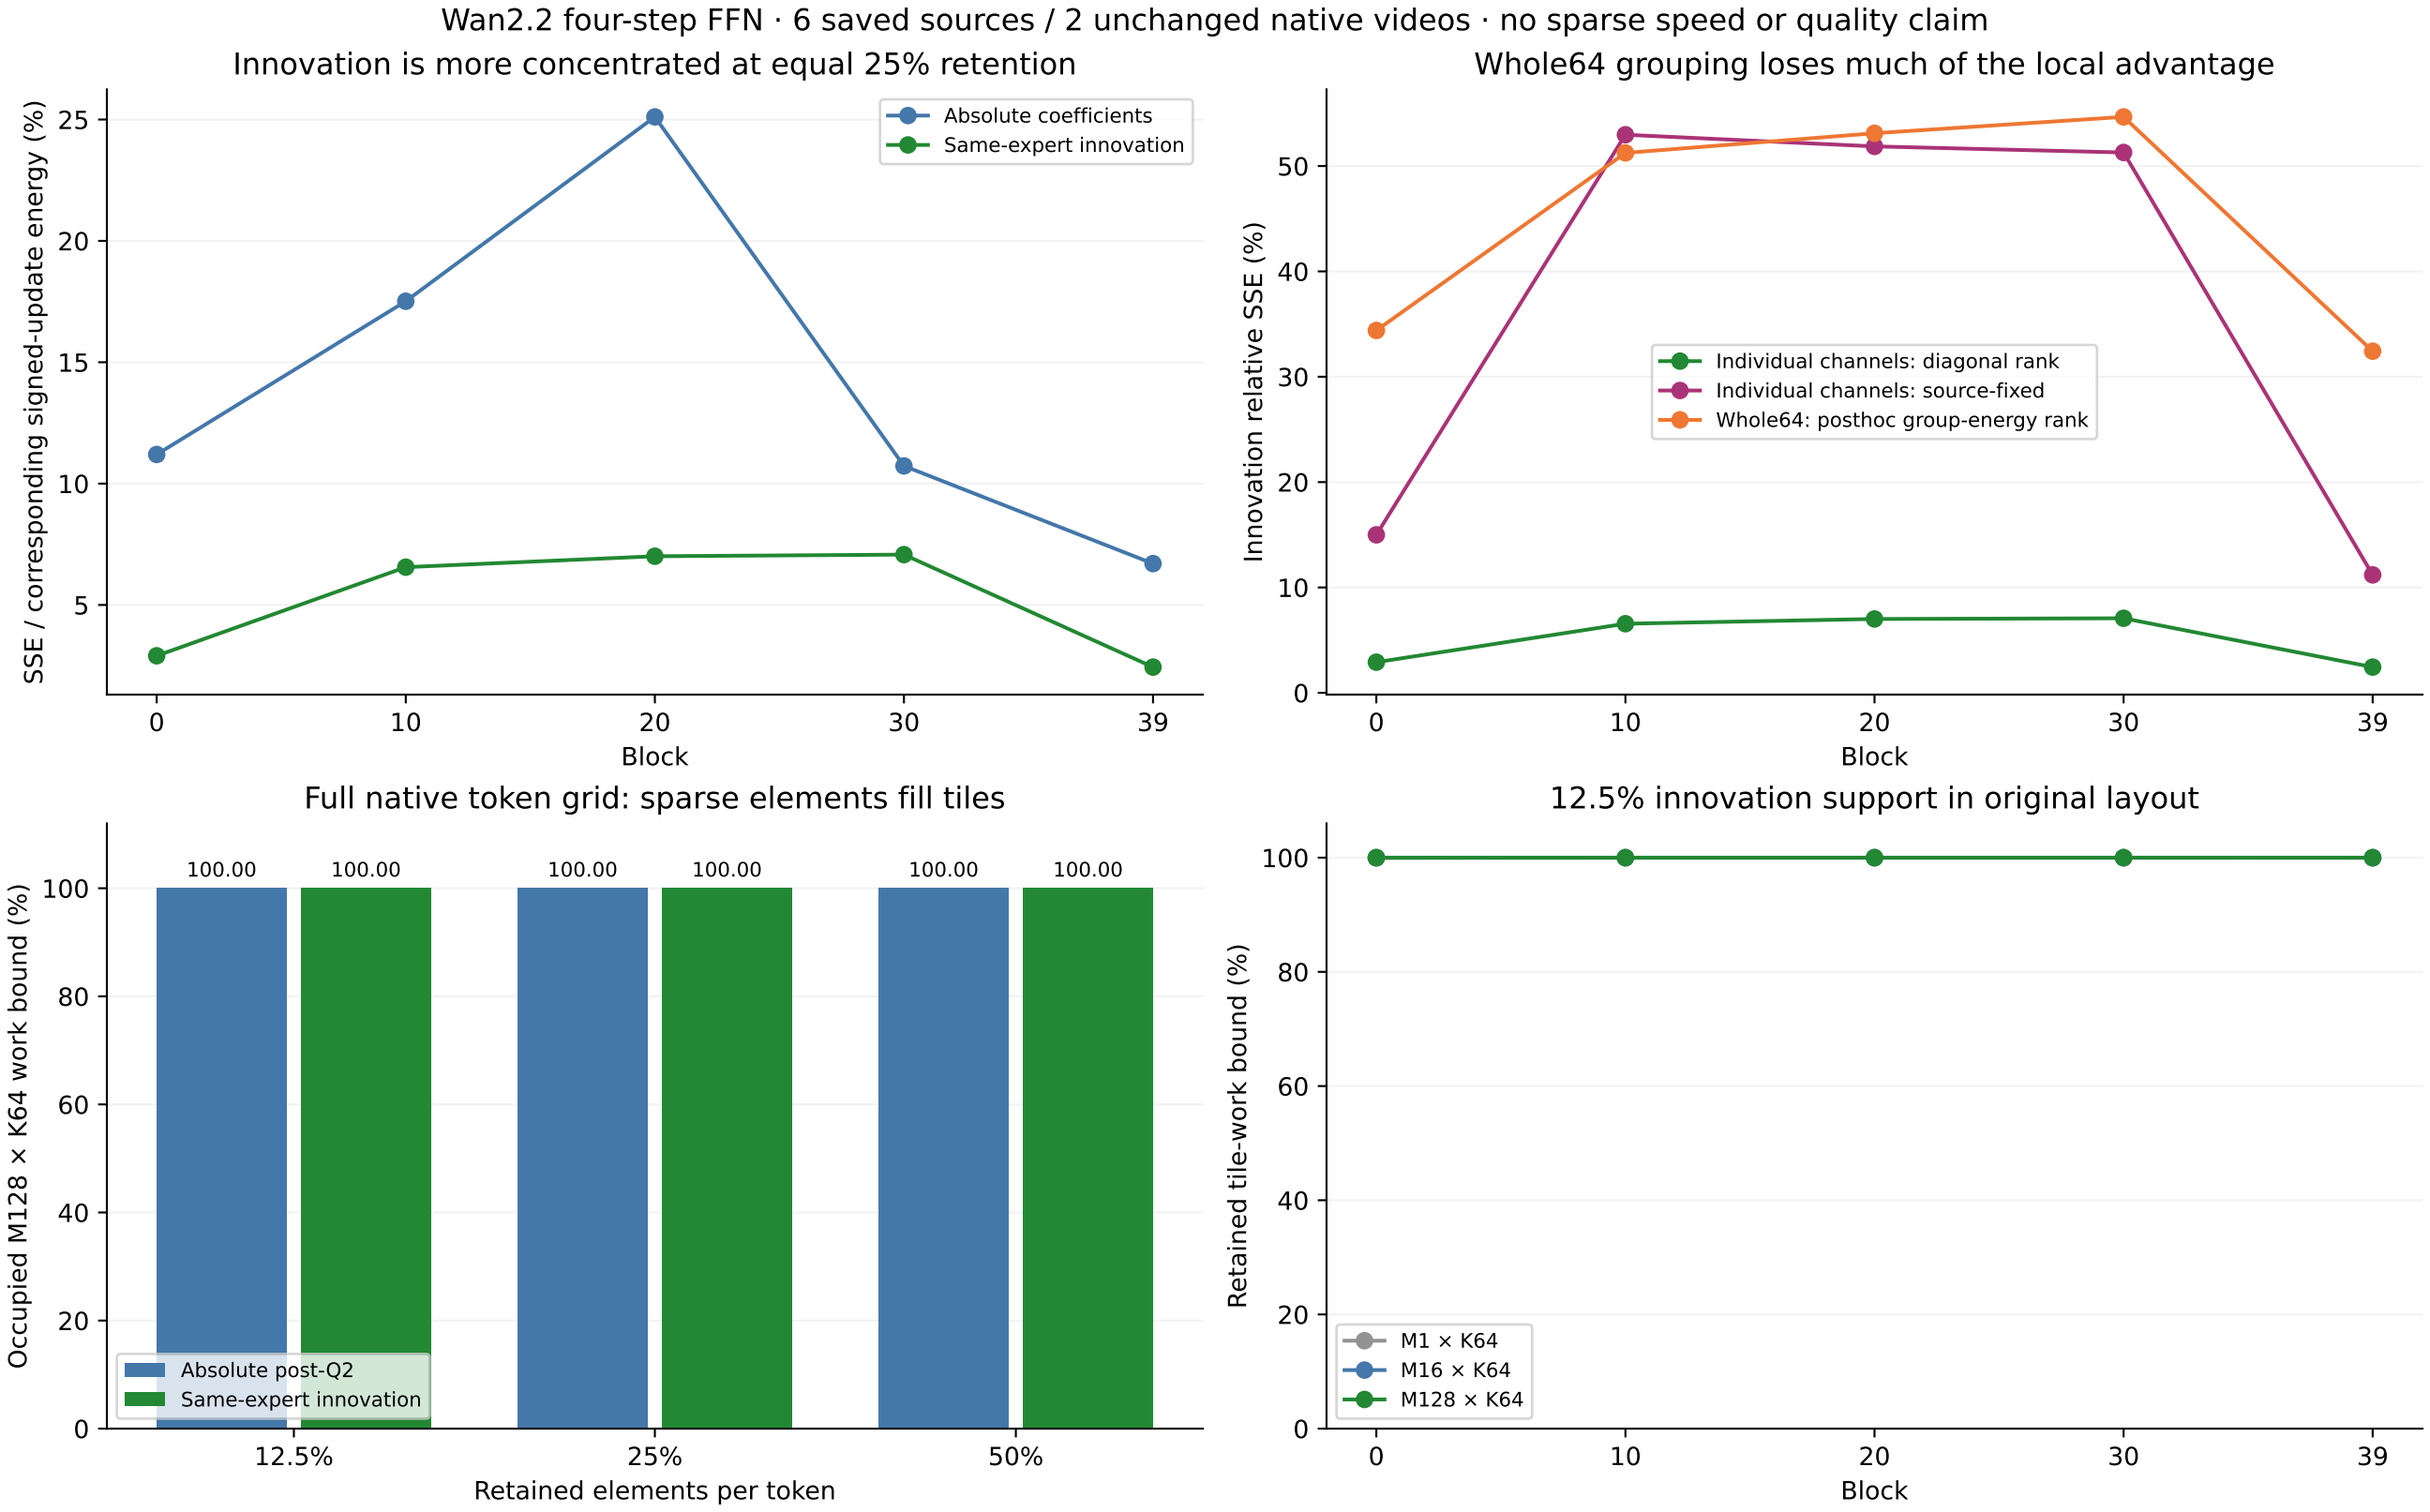
Wan2.2 four-step FFN · 6 saved sources / 2 unchanged native videos · no sparse speed or quality claim
Innovation is more concentrated at equal 25% retention
SSE / corresponding signed-update energy (%)
Absolute coefficients
Same-expert innovation
Block
Whole64 grouping loses much of the local advantage
Innovation relative SSE (%)
Individual channels: diagonal rank
Individual channels: source-fixed
Whole64: posthoc group-energy rank
Block
Full native token grid: sparse elements fill tiles
Occupied M128 × K64 work bound (%)
100.00
Absolute post-Q2
Same-expert innovation
12.5%
25%
50%
Retained elements per token
12.5% innovation support in original layout
Retained tile-work bound (%)
M1 × K64
M16 × K64
M128 × K64
Block